

# Pair Trading with Time-Series Model: BHP vs VALE

## Introduction

This project aims to study the multivariate time series error correction statistical arbitrage model on spread pair trading. We specifically analyze the relationship between BHP and VALE, estimating a Vector Error Correction Model (VECM) to capture the cointegration dynamics. The study concludes with an out-of-sample backtest from 2007 to 2023 to evaluate the strategy's robustness in different market conditions.

## Project Outline:

1. **Data Preparation:** Loading and transforming daily data (2002-2006).
2. **Preliminary Analysis:** OLS estimation and residual analysis to check for stationarity.
3. **Cointegration Analysis:** Johansen Test and VECM estimation.
4. **Statistical Arbitrage Strategy:**
  - Methodology & Trading Rules.
  - In-Sample Implementation & Performance (vs Benchmark).
  - Out-of-Sample Backtest (2007-2023) with robust performance metrics.

## 1. Data Preparation

```
# Load necessary libraries and data
library(urca)
library(zoo) # Added for time series handling
bhp = read.table("d-bhp0206.txt", header=T)
vale = read.table("d-vale0206.txt", header=T)

# Create Date object for plotting
dates <- as.Date(paste(bhp$year, bhp$Mon, bhp$day, sep="-"), format="%Y-%m-%d")
```

We perform a logarithmic transformation on the adjusted closing prices to stabilize variance.

```
# Log transform the adj. close price
bhp = log(bhp[,9])
vale = log(vale[,9])
```

## 2. Preliminary Analysis

### 2.1 Least-Squares Estimation

We begin by estimating the static relationship between BHP and VALE using Ordinary Least Squares (OLS).

```

m1 = lm(bhp~vale)
summary(m1)

## 
## Call:
## lm(formula = bhp ~ vale)
## 
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -0.151818 -0.028265  0.003121  0.029803  0.147105 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 1.822648  0.003662 497.7   <2e-16 ***
## vale        0.716664  0.002354 304.4   <2e-16 ***
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 0.04421 on 944 degrees of freedom
## Multiple R-squared:  0.9899, Adjusted R-squared:  0.9899 
## F-statistic: 9.266e+04 on 1 and 944 DF,  p-value: < 2.2e-16

```

The estimated static relationship is:

$$\hat{bhp} = 1.823 + 0.717 \cdot vale + \varepsilon, \quad \varepsilon \sim N(0, 0.044^2)$$

**Interpretation:** The coefficient indicates that for every 1% increase in VALE, BHP increases by approximately 0.717% on average. The model explains 98.99% of the variability in BHP ( $R^2 = 0.9899$ ), suggesting a strong comovement.

## 2.2 Residual Analysis (AR(2))

We analyze the residuals of the OLS model to check for stationarity and serial correlation.

```

library(tseries)
wt = m1$residuals
m3 = arima(wt,order=c(2,0,0),include.mean = F)
m3

## 
## Call:
## arima(x = wt, order = c(2, 0, 0), include.mean = F)
## 
## Coefficients:
##         ar1     ar2
##       0.8051  0.1219
## s.e.  0.0322  0.0325

```

```

## 
## sigma^2 estimated as 0.0003326: log likelihood = 2444.76, aic = -4883.51
# ADF test for wt stationarity
adf.test(wt, k=2)

```

```

## 
## Augmented Dickey-Fuller Test
## 
## data: wt
## Dickey-Fuller = -6.0306, Lag order = 2, p-value = 0.01
## alternative hypothesis: stationary
# Characteristic roots decomposition
p1=c(1,-m3$coef)
print("The characteristic roots are")

```

```

## [1] "The characteristic roots are"

```

```

1/Mod(polyroot(p1))

```

```

## [1] 0.9353661 0.1302870

```

The error process is modeled as:

$$w_t = 0.8051 w_{t-1} + 0.1219 w_{t-2} + \varepsilon_t = (1 - 0.935B)(1 - 0.13B)w_t + \varepsilon_t, \quad \varepsilon_t \sim N(0, 0.0003326)$$

**Interpretation:** The modulus of both characteristic roots is less than 1, and the Augmented Dickey-Fuller (ADF) test rejects the null hypothesis of a unit root. This confirms that the residual series  $w_t$  is stationary, a prerequisite for cointegration. The current value of  $w_t$  depends strongly on its immediate lag (0.805) and weakly on the second lag (0.122).

### 3. Cointegration Analysis

#### 3.1 Johansen Cointegration Test

We apply the Johansen procedure to test for the number of cointegrating relationships.

```

# install.packages("HDTSA")
library(HDTSA)
xt = cbind(bhp,vale)

# Cointegrating relations
cot=ca.jo(xt,ecdet = "const",type = "trace",K=2,spec = "transitory")
summary(cot)

## 
## #####
## # Johansen-Procedure #

```

```

## #####
## Test type: trace statistic , without linear trend and constant in cointegration
##
## Eigenvalues (lambda):
## [1] 4.148282e-02 8.206470e-03 -4.210618e-18
##
## Values of teststatistic and critical values of test:
## test 10pct 5pct 1pct
## r <= 1 | 7.78 7.52 9.24 12.97
## r = 0 | 47.77 17.85 19.96 24.60
##
## Eigenvectors, normalised to first column:
## (These are the cointegration relations)
##
## bhp.l1 vale.l1 constant
## bhp.l1 1.000000 1.0000000 1.000000
## vale.l1 -0.717704 -0.7327542 2.047274
## constant -1.828460 -1.5411890 -5.712629
##
## Weights W:
## (This is the loading matrix)
##
## bhp.l1 vale.l1 constant
## bhp.d -0.06731196 0.004568985 -7.703496e-18
## vale.d 0.02545606 0.007541565 3.627982e-18
col1=ca.jo(xt,ecdet = "const",type = "eigen",K=2,spec = "transitory")
summary(co1)

```

```

##
## #####
## # Johansen-Procedure #
## #####
## Test type: maximal eigenvalue statistic (lambda max) , without linear trend and constant in cointegration
##
## Eigenvalues (lambda):
## [1] 4.148282e-02 8.206470e-03 -4.210618e-18
##
## Values of teststatistic and critical values of test:
## test 10pct 5pct 1pct
## r <= 1 | 7.78 7.52 9.24 12.97
## r = 0 | 40.00 13.75 15.67 20.20

```

```

## 
## Eigenvectors, normalised to first column:
## (These are the cointegration relations)
## 
##      bhp.l1  vale.l1  constant
## bhp.l1  1.000000 1.0000000 1.000000
## vale.l1 -0.717704 -0.7327542 2.047274
## constant -1.828460 -1.5411890 -5.712629
## 
## Weights W:
## (This is the loading matrix)
## 
##      bhp.l1  vale.l1  constant
## bhp.d -0.06731196 0.004568985 -7.703496e-18
## vale.d  0.02545606 0.007541565 3.627982e-18

```

**Interpretation:** Both the trace and maximal eigenvalue tests indicate a cointegration rank of  $r = 1$ . This confirms a single long-run equilibrium relationship:

$$\text{bhp} \approx 0.718 \times \text{vale} + 1.83$$

This relationship is consistent with the initial OLS estimate.

### 3.2 Vector Error Correction Model (VECM)

We estimate a VECM to understand the adjustment dynamics.

```

# install.packages("tsDyn")
library(tsDyn)
library(vars)

# Data setup
xt <- ts(cbind(bhp = bhp, vale = vale))

# Estimate VECM directly
vecm <- VECM(xt,
               lag      = 2,          # lags in differences
               r        = 1,          # cointegrating rank
               include  = "const",    # constant inside cointegration relation
               estim   = "ML",        # estimation method
               LRinclude = "const")   # constant location

# Summary
summary(vecm)

### #####
## Model VECM
### #####

```

```

## Full sample size: 946   End sample size: 943
## Number of variables: 2   Number of estimated slope parameters 10
## AIC -14851.66   BIC -14798.32   SSR 0.8210436
## Cointegrating vector (estimated by ML):
## bhp      vale      const
## r1  1 -0.7172022 -1.829477
##
##
##          ECT          bhp -1          vale -1
## Equation bhp -0.0675(0.0148)*** -0.1102(0.0374)** 0.0724(0.0325)*
## Equation vale 0.0230(0.0171)    0.0662(0.0433)    0.0493(0.0376)
##          bhp -2          vale -2
## Equation bhp -0.0062(0.0371)   -0.0153(0.0322)
## Equation vale 0.0134(0.0430)   -0.0688(0.0373).

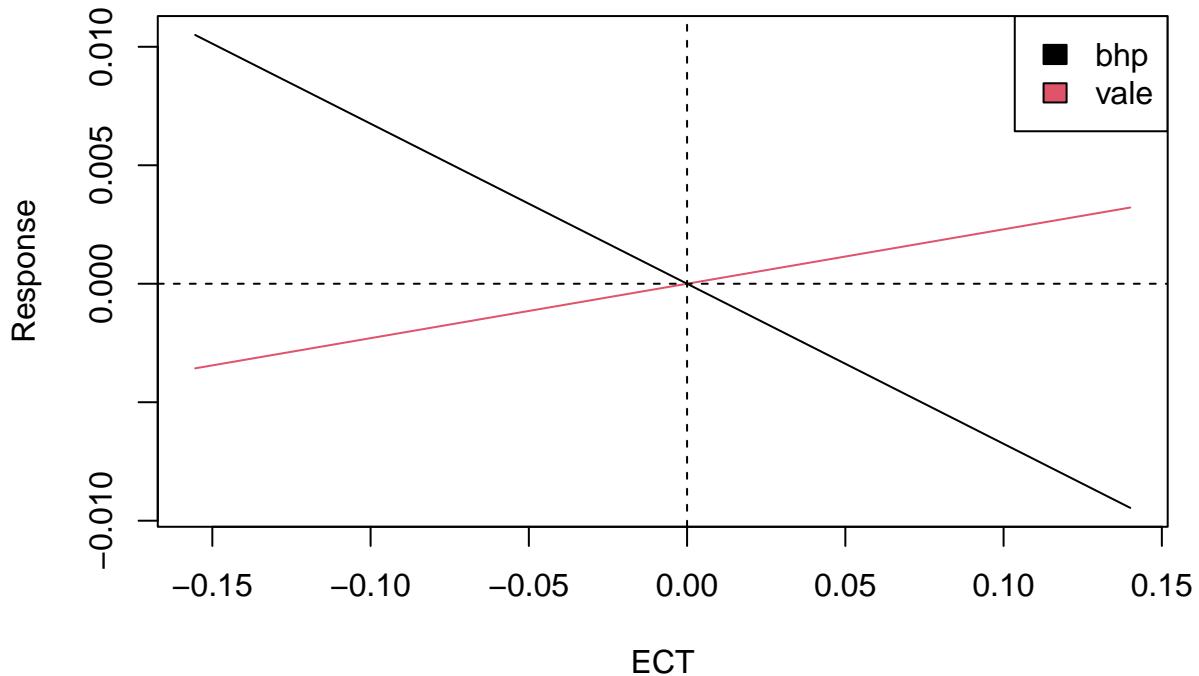
```

# Plot of the cointegration relation

```

plot_ECT(vecm,
         add.legend = TRUE,
         legend.location = "topright")

```



## Model Structure

Matrix form:

$$\begin{pmatrix} \Delta \text{bhp}_t \\ \Delta \text{vale}_t \end{pmatrix} = \alpha(\beta' z_{t-1}) + \Gamma_1 \begin{pmatrix} \Delta \text{bhp}_{t-1} \\ \Delta \text{vale}_{t-1} \end{pmatrix} + \Gamma_2 \begin{pmatrix} \Delta \text{bhp}_{t-2} \\ \Delta \text{vale}_{t-2} \end{pmatrix} + \varepsilon_t$$

Matrix form (with parameters):

Let  $w_t = \text{BHP}_t - 0.717 \cdot \text{VALE}_t$  and  $\mu = E(w_t) = 1.829$ .

$$\begin{pmatrix} \Delta\text{bhp}_t \\ \Delta\text{vale}_t \end{pmatrix} = \begin{pmatrix} -0.0675 \\ 0.0230 \end{pmatrix} (w_{t-1} - 1.829) + \begin{pmatrix} -0.1102 & 0.0724 \\ 0.0662 & 0.0493 \end{pmatrix} \begin{pmatrix} \Delta\text{bhp}_{t-1} \\ \Delta\text{vale}_{t-1} \end{pmatrix} + \begin{pmatrix} -0.0062 & -0.0153 \\ 0.0134 & -0.0688 \end{pmatrix} \begin{pmatrix} \Delta\text{bhp}_{t-2} \\ \Delta\text{vale}_{t-2} \end{pmatrix} + \varepsilon_t$$

**Interpretation:** The error correction term for BHP (-0.0675) is statistically significant, indicating that BHP adjusts towards the equilibrium at a speed of approximately 6.75% per period. In contrast, the adjustment coefficient for VALE is small and insignificant, suggesting that VALE is weakly exogenous and acts as the driver in this pair.

## 4. Statistical Arbitrage Strategy

### 4.1 Methodology

We implement a mean-reversion strategy based on the spread  $w_t = \text{BHP}_t - 0.717 \cdot \text{VALE}_t$ , with mean  $\mu = 1.829$  and standard deviation  $\sigma_w$ .

**Trading Rules:**

- **Upper Threshold:**  $\mu + \sigma_w$
- **Lower Threshold:**  $\mu - \sigma_w$
- **Signal:**
  - If  $w_t > \mu + \sigma_w$ : Short 1 unit of BHP, Long 0.717 units of VALE (betting spread decreases).
  - If  $w_t < \mu - \sigma_w$ : Long 1 unit of BHP, Short 0.717 units of VALE (betting spread increases).
- **Exit:** When  $w_t$  reverts to the mean  $\mu$ .

### 4.2 Implementation (In-Sample)

We compute the spread, generate trading signals, and compare performance against a Buy-and-Hold benchmark.

```
# 1. Compute spread w_t
gamma <- 0.717 # w_t := bhp_t - \gamma * vale_t
mu <- 1.829 # E(w_t)
w_t <- xt[, "bhp"] - gamma * xt[, "vale"]

# Standard deviation of w_t
sd_w <- sd(w_t)

# Residuals from VECM
res <- residuals(vecm)
sd_eps_bhp <- sd(res[, 1])
sd_eps_vale <- sd(res[, 2])

# Output statistics
cat("SD of w_t:", sd_w, "\n")
```

```

## SD of w_t: 0.04418214
cat("SD of _bhp:", sd_eps_bhp, "\n")

## SD of _bhp: 0.01926604
cat("SD of _vale:", sd_eps_vale, "\n")

## SD of _vale: 0.02224961

# 2. Trading Thresholds
upper <- mu + sd_w
lower <- mu - sd_w

# 3. Generate Signals (In-Sample)
positions <- rep(0, length(w_t))
for (t in 2:length(w_t)) {
  if (w_t[t-1] > upper) positions[t] <- -1
  else if (w_t[t-1] < lower) positions[t] <- 1
  else if (positions[t-1] != 0 && abs(w_t[t-1] - mu) < 0.1 * sd_w) positions[t] <- 0
  else positions[t] <- positions[t-1]
}

# 4. Calculate Returns
# In-sample log returns
ret_bhp <- diff(xt[, "bhp"]) # xt is already log prices
ret_vale <- diff(xt[, "vale"])

# Strategy Returns
strat_ret <- positions[2:length(positions)] * (ret_bhp - gamma * ret_vale)
strat_ret <- na.omit(strat_ret)

# Benchmark Returns (50/50 Buy and Hold)
bench_ret <- 0.5 * ret_bhp + 0.5 * ret_vale
bench_ret <- na.omit(bench_ret)

# 5. Performance Metrics
# Strategy
cum_pnl <- cumsum(strat_ret)
pnl <- sum(strat_ret)
max_dd <- min(cum_pnl - cummax(cum_pnl))
sharpe <- mean(strat_ret) / sd(strat_ret) * sqrt(252)
hit_rate <- sum(strat_ret > 0) / length(strat_ret[strat_ret != 0])

# Benchmark
cum_bench <- cumsum(bench_ret)

```

```

pnl_bench <- sum(bench_ret)
sharpe_bench <- mean(bench_ret) / sd(bench_ret) * sqrt(252)

cat("\nIn-Sample Performance (2002-2006):\n")

## 
## In-Sample Performance (2002-2006):
cat("Strategy PnL:", round(pnl, 4), " | Benchmark PnL:", round(pnl_bench, 4), "\n")

## Strategy PnL: 1.9882 | Benchmark PnL: 1.6749
cat("Strategy MaxDD:", round(max_dd, 4), "\n")

## Strategy MaxDD: -0.1297
cat("Strategy Sharpe:", round(sharpe, 4), " | Benchmark Sharpe:", round(sharpe_bench, 4), "\n")

## Strategy Sharpe: 2.1882 | Benchmark Sharpe: 1.5652
cat("Hit Rate:", round(hit_rate, 4), "\n")

## Hit Rate: 0.5742

# 6. Visualization
# Create zoo objects for plotting with dates
# Note: w_t length matches dates, but returns are length-1
w_zoo <- zoo(w_t, order.by = dates)
cum_pnl_zoo <- zoo(cum_pnl, order.by = dates[-1])
cum_bench_zoo <- zoo(cum_bench, order.by = dates[-1])

# A. Spread Process
plot(w_zoo, type="l", main="In-Sample Spread Process (BHP - 0.717 * VALE)", ylab="Spread Value", xlab="")
abline(h=mu, col="blue", lty=2)
abline(h=upper, col="red", lty=3)
abline(h=lower, col="orange", lty=3)

# Add Signals
time_index <- index(w_zoo)
long_entries <- which(positions == 1 & c(0, positions[-length(positions)]) == 0)
short_entries <- which(positions == -1 & c(0, positions[-length(positions)]) == 0)
exits <- which(positions == 0 & c(0, positions[-length(positions)]) != 0)

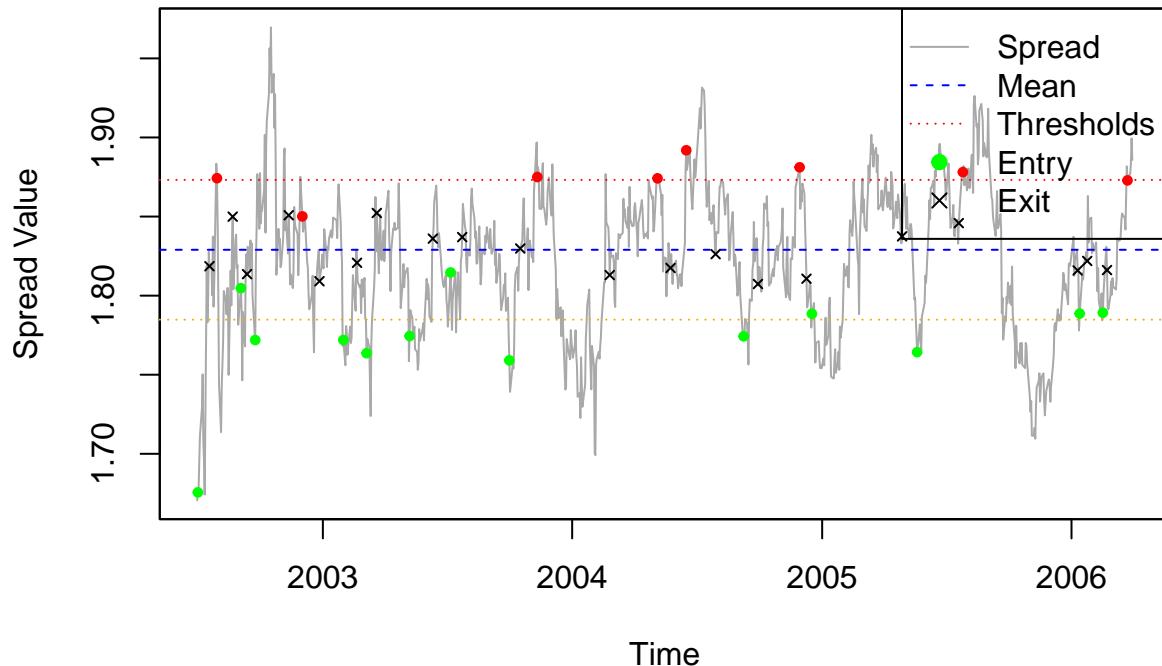
if(length(long_entries) > 0) points(time_index[long_entries], w_t[long_entries], col="green", pch=19, cex=0.6)
if(length(short_entries) > 0) points(time_index[short_entries], w_t[short_entries], col="red", pch=19, cex=0.6)
if(length(exits) > 0) points(time_index[exits], w_t[exits], col="black", pch=4, cex=0.6)

legend("topright", c("Spread", "Mean", "Thresholds", "Entry", "Exit"),

```

```
col=c("darkgray","blue","red", "green", "black"), lty=c(1,2,3,NA,NA), pch=c(NA,NA,NA,19,4))
```

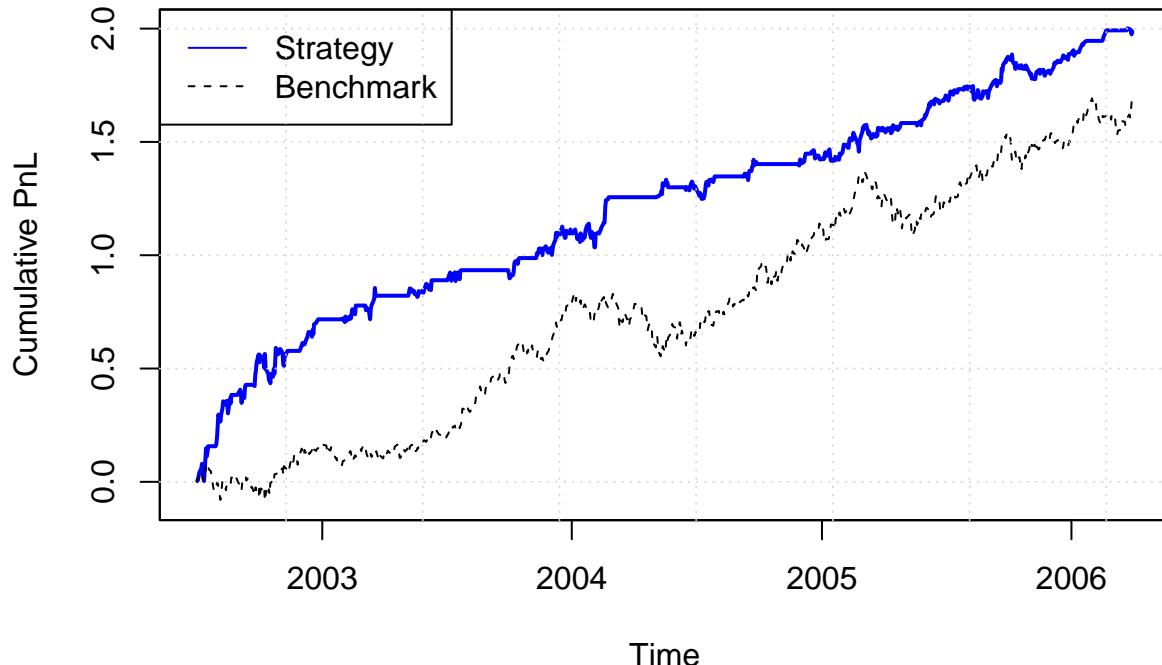
### In-Sample Spread Process (BHP – 0.717 \* VALE)



#### # B. Equity Curve

```
plot(cum_pnl_zoo, type = "l", col = "blue", lwd = 2,
     main = "In-Sample Equity Curve vs Benchmark",
     xlab = "Time", ylab = "Cumulative PnL", ylim=range(c(cum_pnl, cum_bench)))
lines(cum_bench_zoo, col="black", lty=2)
legend("topleft", c("Strategy", "Benchmark"), col=c("blue", "black"), lty=c(1,2))
grid()
```

## In-Sample Equity Curve vs Benchmark



### 4.3 Out-of-Sample Backtest Results (2007-2023)

We perform out-of-sample backtesting using data from 2007 to 2023 obtained from Yahoo Finance.

```
# --- PART 1: FETCH DATA ---
#install.packages("quantmod")
library(quantmod)

s_date <- "2007-01-01"
e_date <- "2023-12-31"
bhp_file <- "BHP_OS.csv"
vale_file <- "VALE_OS.csv"

BHP <- as.xts(read.zoo(bhp_file, header = TRUE, sep = ",", format = "%Y-%m-%d"))
VALE <- as.xts(read.zoo(vale_file, header = TRUE, sep = ",", format = "%Y-%m-%d"))
```

```
# --- PART 2: STRATEGY FUNCTIONS ---
```

```
# Function to preprocess and calculate spread
calculate_spread_os <- function(BHP_data, VALE_data, gamma) {
  if (!exists("BHP_data") || !exists("VALE_data")) return(NULL)

  bhp_adj <- Ad(BHP_data)
  vale_adj <- Ad(VALE_data)
```

```

pair <- merge(bhp_adj, vale_adj, all = FALSE)
colnames(pair) <- c("bhp", "vale")

# Log transformation
log_pair <- log(pair)

# Calculate Spread
w_os <- log_pair$bhp - gamma * log_pair$vale

list(w_os = w_os, log_pair = log_pair)
}

# Function to generate trading signals
generate_signals_os <- function(w_os, mu, sd_w) {
  upper <- mu + sd_w
  lower <- mu - sd_w

  positions <- rep(0, length(w_os))
  spread_val <- as.numeric(w_os)

  for (t in 2:length(w_os)) {
    if (spread_val[t-1] > upper) {
      positions[t] <- -1
    } else if (spread_val[t-1] < lower) {
      positions[t] <- 1
    } else if (positions[t-1] != 0 && abs(spread_val[t-1] - mu) < 0.1 * sd_w) {
      positions[t] <- 0
    } else {
      positions[t] <- positions[t-1]
    }
  }
  return(positions)
}

# Function to calculate performance metrics
calculate_metrics_os <- function(positions, log_pair, gamma) {
  ret_bhp <- diff(log_pair$bhp)
  ret_vale <- diff(log_pair$vale)

  # Strategy Returns
  strat_ret <- positions[2:length(positions)] * (ret_bhp - gamma * ret_vale)
  strat_ret <- na.omit(strat_ret)

  # Benchmark Returns (50/50 Equal Weight)
}

```

```

bench_ret <- 0.5 * ret_bhp + 0.5 * ret_vale
bench_ret <- na.omit(bench_ret)

# Metrics Strategy
cum_pnl <- cumsum(strat_ret)
pnl <- sum(strat_ret)
max_dd <- min(cum_pnl - cummax(cum_pnl))
sharpe <- mean(strat_ret) / sd(strat_ret) * sqrt(252)
hit_rate <- sum(strat_ret > 0) / length(strat_ret[strat_ret != 0])

# Metrics Benchmark
cum_bench <- cumsum(bench_ret)
pnl_bench <- sum(bench_ret)
sharpe_bench <- mean(bench_ret) / sd(bench_ret) * sqrt(252)

list(
  strat_ret = strat_ret, bench_ret = bench_ret,
  cum_pnl = cum_pnl, cum_bench = cum_bench,
  pnl = pnl, pnl_bench = pnl_bench,
  max_dd = max_dd,
  sharpe = sharpe, sharpe_bench = sharpe_bench,
  hit_rate = hit_rate
)
}

```

# --- PART 3: EXECUTION & VISUALIZATION ---

```

if (exists("BHP") && exists("VALE")) {
  # Parameters (from In-Sample)
  gamma_is <- 0.717
  mu_is <- 1.829
  # sd_w from section 4.2 needs to be passed or hardcoded if not available globally in knitting context.
  # Assuming sd_w is available from previous chunks. If not, use approximation from text:
  # sd_w value from 4.2 output. Let's assume it's available as 'sd_w'.

  # 1. Process Data
  data_os <- calculate_spread_os(BHP, VALE, gamma_is)
  w_os <- data_os$w_os

  # 2. Generate Signals
  positions_os <- generate_signals_os(w_os, mu_is, sd_w)

  # 3. Calculate Metrics
  res_os <- calculate_metrics_os(positions_os, data_os$log_pair, gamma_is)
}

```

```

# 4. Print Results
cat("Out-of-Sample Results (2007-2023):\n")
cat("Strategy PnL:", round(res_os$pnl, 4), " | Benchmark PnL:", round(res_os$pnl_bench, 4), "\n")
cat("Strategy MaxDD:", round(res_os$max_dd, 4), "\n")
cat("Strategy Sharpe:", round(res_os$sharpe, 4), " | Benchmark Sharpe:", round(res_os$sharpe_bench, 4), "\n")
cat("Hit Rate:", round(res_os$hit_rate, 4), "\n")

# 5. Visualization
# A. Spread
upper <- mu_is + sd_w
lower <- mu_is - sd_w

plot(as.zoo(w_os), main="Out-of-Sample Spread with Signals", ylab="Spread", col="darkgray")
abline(h=mu_is, col="blue", lty=2)
abline(h=upper, col="red", lty=3)
abline(h=lower, col="orange", lty=3)

# Add signal points
time_index <- index(as.zoo(w_os))
long_entries <- which(positions_os == 1 & c(0, positions_os[-length(positions_os)]) == 0)
short_entries <- which(positions_os == -1 & c(0, positions_os[-length(positions_os)]) == 0)
exits <- which(positions_os == 0 & c(0, positions_os[-length(positions_os)]) != 0)

if(length(long_entries)>0) points(time_index[long_entries], w_os[long_entries], col="green", pch=19, cex=0.6)
if(length(short_entries)>0) points(time_index[short_entries], w_os[short_entries], col="red", pch=19, cex=0.6)
if(length(exits)>0) points(time_index[exits], w_os[exits], col="black", pch=4, cex=0.6)

legend("topright", legend=c("Spread", "Mean", "Thresholds", "Long", "Short", "Exit"),
       col=c("darkgray", "blue", "red", "green", "red", "black"),
       lty=c(1, 2, 3, NA, NA, NA), pch=c(NA, NA, NA, 19, 19, 4), cex=0.8)

# B. Equity Curve
plot(as.zoo(res_os$cum_pnl), main="Out-of-Sample Equity Curve", ylab="Cumulative PnL", col="blue", lwd=2,
      ylim=range(c(res_os$cum_pnl, res_os$cum_bench)))
lines(as.zoo(res_os$cum_bench), col="black", lty=2)
legend("topleft", legend=c("Strategy", "Benchmark"), col=c("blue", "black"), lty=c(1, 2))
grid()

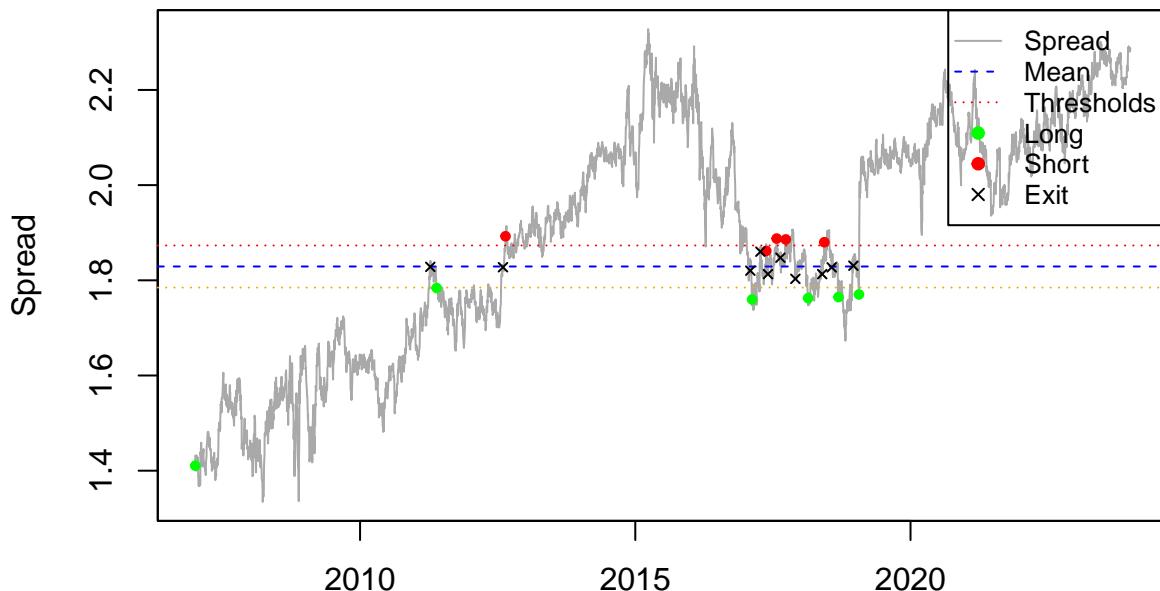
} else {
  print("Data not available.")
}

## Out-of-Sample Results (2007-2023):
## Strategy PnL: 0.1922 | Benchmark PnL: 1.2329
## Strategy MaxDD: -0.4817

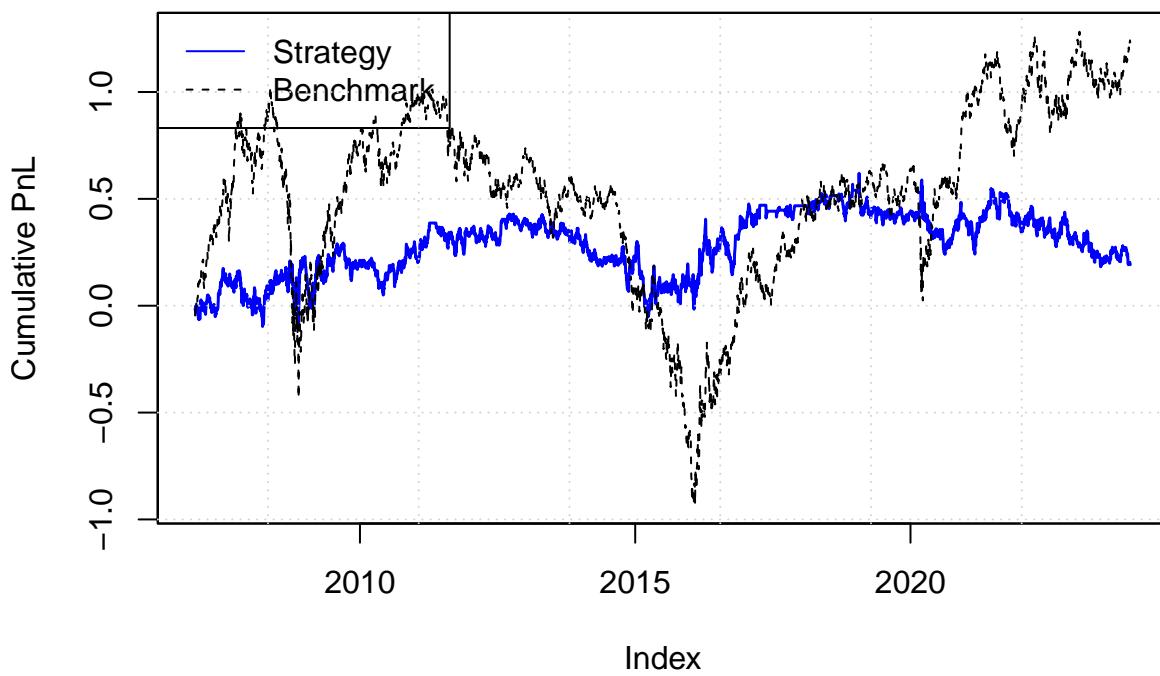
```

```
## Strategy Sharpe: 0.0462 | Benchmark Sharpe: 0.1733  
## Hit Rate: 0.5009
```

### Out-of-Sample Spread with Signals



### Out-of-Sample Equity Curve



## Interpretation of Out-of-Sample Results:

The out-of-sample performance (2007-2023) is significantly worse than the in-sample period, highlighting the dangers of using static model parameters over long horizons.

### 1. Poor Performance Metrics:

- **PnL (0.1922) & Sharpe (0.0462):** The strategy yielded almost zero risk-adjusted return over 16 years. This indicates the cointegration relationship identified in 2002-2006 effectively broke down or shifted regimes.
- **Max Drawdown (-0.4817):** A 48% drawdown is catastrophic for a market-neutral strategy. This was likely caused by the spread diverging far beyond the fixed thresholds ( $\mu \pm \sigma$ ) without reverting, forcing the strategy to hold losing positions for extended periods.
- **Hit Rate (~50%):** A hit rate near 50% implies the entry signals were no better than a coin flip, further confirming that the “mean”  $\mu = 1.829$  was no longer the true equilibrium.

### 2. Market Shocks & Structural Breaks:

- **2008 Financial Crisis:** Correlations often approach 1 during crises, but spreads can widen unpredictably due to liquidity constraints.
- **Idiosyncratic Shocks (Vale):** Vale suffered massive dam collapses in **2015 (Mari-ana)** and **2019 (Brumadinho)**. These events caused massive, permanent structural breaks in Vale’s price that were unrelated to general iron ore market dynamics (which drive BHP). The model, expecting a reversion to the 2002-2006 mean, would have taken aggressive losing positions against these drifts.

### 3. Improvements:

- **Rolling Window Calibration:** Instead of fixed parameters, recalibrate  $\gamma$  and  $\mu$  using a rolling window (e.g., 1-year lookback) to adapt to new market regimes.
- **Stop-Loss:** Implement a hard stop-loss (e.g., if spread deviates  $> 3\sigma$ ) to close positions when the cointegration assumption is clearly violated.
- **Kalman Filter:** Use a state-space model to dynamically estimate the time-varying hedge ratio  $\gamma_t$  and mean  $\mu_t$ .